

SensiSpace

Satellite Land Cover Classification

A Deep Learning Approach using Baby Dragon Hatchling Architecture
vs ResNet-18 on EuroSAT Dataset (27,000 images)

ISRO RESPOND Initiative - Deep Learning Research
Major Project Report

1. Problem Statement

Satellite remote sensing generates vast amounts of high-resolution imagery used for monitoring Earth's surface, environmental changes, and natural resources. Organizations like ISRO (Indian Space Research Organisation) collect this data, but manually analyzing it is impractical at scale.

This project introduces an intelligent system that automatically classifies satellite images into 10 land cover types using deep learning, and visually highlights which parts of the image correspond to which land type.

We compare a standard industry baseline (ResNet-18) against a novel, biologically-inspired architecture called Baby Dragon Hatchling (BDH) to demonstrate that purpose-built architectures can outperform generic ones with fewer parameters.

2. Dataset - EuroSAT

EuroSAT is a benchmark dataset of 27,000 georeferenced satellite images from the Sentinel-2 satellite (European Space Agency). Each image is 64x64 pixels in RGB and labeled with one of 10 land cover classes.

Dataset Properties

Property	Value
Source	Sentinel-2 satellite
Total Images	27,000
Image Size	64 x 64 pixels, RGB
Classes	10 land cover types
Train/Val Split	80% / 20%

The 10 Land Cover Classes

Class	Description
AnnualCrop	Wheat, corn, rice fields
Forest	Dense tree cover (deciduous/coniferous)
HerbaceousVegetation	Grasslands, meadows, scrubland
Highway	Major roads and expressways
Industrial	Factories, warehouses
Pasture	Grazing land for livestock
PermanentCrop	Orchards, vineyards
Residential	Houses, apartments, neighborhoods
River	Rivers and streams
SeaLake	Seas, lakes, reservoirs

3. Model 1 - ResNet-18 (CNN Baseline)

What is ResNet?

ResNet (Residual Network) is a standard Convolutional Neural Network developed by Microsoft Research in 2015. The key innovation is skip connections (residual connections) that allow gradients to flow directly through the network, solving the vanishing gradient problem in deep networks. ResNet-18 has 18 layers and approximately 11.2 million parameters.

Architecture Pipeline

```
Input (3x64x64)
| -> Conv 7x7, stride 2 + BatchNorm + ReLU + MaxPool 3x3
|   -> 64 channels
| -> Layer 1: 2x BasicBlock (64 channels)
| -> Layer 2: 2x BasicBlock (128 channels, stride 2)
| -> Layer 3: 2x BasicBlock (256 channels, stride 2)
| -> Layer 4: 2x BasicBlock (512 channels, stride 2)
| -> Global Average Pooling (512x1x1)
| -> Fully Connected: 512 -> 10 classes
```

BasicBlock (Building Block)

```
Input ----> Conv3x3 -> BN -> ReLU -> Conv3x3 -> BN ----> ADD -> ReLU -> Output
|                                                                 |
+----- Skip Connection (identity) -----+
```

The skip connection adds the original input directly to the output. The block only needs to learn the difference (residual) between input and desired output, making training much easier. This is what allows ResNet to train successfully with many layers.

Transfer Learning

We use ResNet-18 pretrained on ImageNet (1.3 million natural images). The model already knows basic visual features like edges, textures, and shapes. We replace the final fully connected layer (512->1000) with (512->10) and fine-tune the entire network on EuroSAT satellite imagery. This approach is called transfer learning - leveraging knowledge from one domain (natural images) for another (satellite imagery).

4. Model 2 - Baby Dragon Hatchling (BDH)

Baby Dragon Hatchling is a novel, biologically-inspired deep learning architecture designed specifically for satellite remote sensing image analysis. Developed under the ISRO RESPOND initiative, it draws from neuroscience - how the human visual cortex processes images - and incorporates four specialized modules that standard CNNs lack.

Architecture Pipeline

```

Input (3x64x64)
|-> Stem: Conv3x3 -> BN -> ReLU -> Conv3x3 -> BN -> ReLU (3 -> 64 ch)
|
|-> BDH Stage 1 (64 ch, 64x64) -> MaxPool 2x2 -> 32x32
|-> BDH Stage 2 (128 ch, 32x32) -> MaxPool 2x2 -> 16x16
|-> BDH Stage 3 (256 ch, 16x16) -> MaxPool 2x2 -> 8x8
|
|-> Recurrent Refinement (2 iterations)
|
|-> Global Average Pool (256x1x1)
|-> Dropout(0.3) -> FC(256->128) -> ReLU
|-> Dropout(0.2) -> FC(128->10)

```

Inside Each BDH Stage

```

Input --> Multi-Scale Pyramid --> Lateral Inhibition --> Contextual Attention --> ADD --> ReLU
|                                                                                               |
+----- 1x1 Conv Residual Projection -----+

```

Each stage chains three biologically-inspired modules with a residual connection. The residual projection (1x1 convolution) adjusts channel dimensions when they change between stages.

Module 1: Multi-Scale Feature Pyramid

Problem: Satellite images contain objects at many different sizes - a small house vs. a large forest. A single 3x3 convolution can only "see" a small receptive field.

Solution: Process the input through three parallel convolutions with different kernel sizes simultaneously, then concatenate the results:

```
Input --+--> Conv 3x3 + BN + ReLU (fine detail: edges, small objects)    --+
          +--> Conv 5x5 + BN + ReLU (medium features: buildings, patches) --+--> Concatenate
          +--> Conv 7x7 + BN + ReLU (large patterns: forests, water bodies) --+
```

The 3x3 branch captures fine edges (road boundaries), the 5x5 captures medium structures (buildings), and the 7x7 captures broad patterns (forest canopy). By concatenating all three, the model gets a rich, multi-resolution representation at every stage. This is critical for EuroSAT where land cover objects vary dramatically in scale.

Module 2: Lateral Inhibition Block

Biological inspiration: In the human retina, when a neuron fires, it suppresses its neighboring neurons. This is called lateral inhibition - it is why we perceive sharper edges between light and dark areas than actually exist in the image.

```
surround = depthwise_conv_3x3(x)      # Estimate surrounding activity
surround = batch_norm(surround)
output   = x - alpha * surround        # Subtract surround (alpha is learnable)
output   = ReLU(output)
```

The depthwise convolution estimates the "surround" response for each feature channel. Subtracting it amplifies differences between adjacent regions. The learnable parameter alpha (initialized to 0.3) controls inhibition strength. Result: sharper feature boundaries - critical for distinguishing adjacent land types (e.g., where a river meets a forest).

Module 3: Contextual Attention

Problem: Not all features are equally important. In a satellite image, water texture might be more informative than road edges for classifying "River."

Solution: A dual attention mechanism that learns to focus on what matters:

Channel Attention (WHAT features matter):

```
Global Avg Pool -> FC(C -> C/8) -> ReLU -> FC(C/8 -> C) -> Sigmoid
Output: a weight per channel, e.g., "90% attention to water features, 10% to edges"
```

Spatial Attention (WHERE to look):

```
AvgPool(channels) + MaxPool(channels) -> Concat -> Conv 7x7 -> Sigmoid
Output: a weight per pixel, e.g., "focus on center-right where the river is"
```

The input is first reweighted by channel importance, then by spatial importance. This gives the model selective focus - just like how humans look at a complex image and immediately focus on the relevant parts.

Module 4: Recurrent Refinement

Biological inspiration: The brain does not process visual information in a single pass. There are feedback connections from higher visual areas back to lower ones - the brain iteratively refines its interpretation of what it sees.

```
for iteration in [1, 2]:
    refined = Conv3x3 + BN + ReLU(x)      # Propose a refinement
    gate     = Sigmoid(Conv1x1(refined))  # How much to update? (0 to 1)
    x        = x * gate + refined * (1-gate) # Gated blend
```

Each iteration, the network computes a "refined" version of the features. A sigmoid gate decides how much of the refinement to accept. Where gate is close to 1, the original features are kept (already good). Where gate is close to 0, they are replaced with the refined version (needs improvement). Result: features become progressively cleaner and more discriminative over 2 iterations.

5. Training Pipeline

Both models are trained on the same 80/20 split of EuroSAT with identical augmentation and optimization settings for fair comparison.

Training Configuration

Setting	Value
Optimizer	Adam
Learning Rate	0.001
LR Scheduler	StepLR (x0.1 every 5 epochs)
Loss Function	CrossEntropyLoss
Batch Size	64
Epochs	15
Augmentation	Flip, Rotate +/-15°, Color Jitter
Normalization	ImageNet mean/std

Results

Metric	ResNet-18	BDH
Final Accuracy	81.98%	95.87%
Parameters	~11.2M	~1.8M
Architecture	Transfer Learning	Trained from scratch

Key insight: BDH achieves 14% higher accuracy with 6x fewer parameters - demonstrating that domain-specific architecture design beats brute-force parameter scaling. The biologically-inspired modules (multi-scale processing, lateral inhibition, attention, and recurrent refinement) are more effective than simply stacking more layers.

6. Spectral Land Cover Segmentation

After classification, we generate a visual overlay showing which parts of the image correspond to which land type. This provides interpretable, visual evidence for the model's classification decision.

Why Not Grad-CAM?

Grad-CAM (Gradient-weighted Class Activation Mapping) is a common visualization technique, but it produces blurry, inaccurate blobs for this use case because: (1) Classification models learn "this whole image = Forest" - they do not learn pixel-level boundaries. (2) The activation maps are very low resolution (8x8 or 2x2 pixels upscaled to 64x64), creating unavoidable artifacts.

Our Approach: Spectral Color Analysis

We use spectral analysis - the same principle used in professional satellite remote sensing with indices like NDVI (Normalized Difference Vegetation Index). The approach analyzes actual pixel colors in HSV and LAB color spaces:

- **Green Leaf Index (GLI):** $(2G - R - B) / (2G + R + B)$ - identifies vegetation
- **HSV Hue Analysis:** Green hue = vegetation, Blue hue = water, Low saturation = urban
- **LAB Color Space:** a-channel warmth for urban detection, b-channel for soil

Segmentation Pipeline

1. Convert satellite image to HSV + LAB color spaces
2. Compute vegetation/water/urban spectral indices per pixel
3. Threshold pixels into categories: vegetation, water, urban, road, soil
4. Map each category to the specific EuroSAT class using model probabilities
5. Draw colored overlays + contour borders on the original image
6. Add class labels with confidence percentages

Result: Pixel-accurate boundaries that follow actual land features (river edges, forest boundaries, road paths) - far more accurate than Grad-CAM.

7. System Architecture

The system is a full-stack web application with a FastAPI backend serving PyTorch models and a premium dark-mode frontend.

Backend (FastAPI + PyTorch)

Endpoint	Method	Returns
/api/predict	POST	Both models predictions + spectral overlay
/api/sample	GET	Random EuroSAT sample image
/api/metrics	GET	Training benchmark data

/api/colors	GET	Color palette for all 10 classes
-------------	-----	----------------------------------

Frontend (HTML/CSS/JS)

- 3-slide tabbed interface: ResNet-18, Baby Dragon Hatchling, Comparison
- Color-coded probability distribution charts for all 10 classes
- Side-by-side spectral overlay comparison
- Premium dark-mode design with glassmorphism and micro-animations
- Upload custom images or load random EuroSAT samples

8. Architecture Comparison Summary

Feature	ResNet-18	BDH
Type	Standard CNN	Bio-inspired CNN
Depth	18 layers	3 stages + refinement
Multi-Scale	No (single 3x3)	Yes (3x3 + 5x5 + 7x7)
Attention	None	Channel + Spatial
Lateral Inhibition	None	Yes (learnable)
Recurrent Feedback	None (feed-forward)	Yes (2 iterations)
Skip Connections	Yes (residual)	Yes (residual)
Parameters	~11.2M	~1.8M
Pretrained	Yes (ImageNet)	No (from scratch)
Final Accuracy	81.98%	95.87%

9. Conclusion

This project demonstrates that biologically-inspired neural network architectures can significantly outperform standard CNNs for satellite remote sensing tasks. The Baby Dragon Hatchling architecture achieves 95.87% accuracy with only 1.8 million parameters - 14 percentage points higher than ResNet-18 with 6x fewer parameters.

The four key innovations - multi-scale feature pyramids, lateral inhibition, contextual attention, and recurrent refinement - each contribute to this improvement by mimicking how the human visual system processes complex spatial information. Combined with spectral color-space segmentation for visualization, the system provides both accurate classification and interpretable, pixel-level land cover mapping.

The system is production-ready for demonstration, with a full web interface allowing users to upload satellite images and receive instant classification results with visual land cover highlighting.